

Stephen Lanna

Assignment 4

NLP

04/19/2023

<https://github.com/beep-sl/Assignment-4>

	P@1	MRR
Part 1	0	0.126
Part 2	0	0.000473
Part 3	0	0.004728

Overall, the pre-trained model worked better for finding similar questions. In my models, testing via MRR was significantly lower. I believe this has to do with the data supplied to the model. The Law Stack Exchange dataset was significantly smaller (I assume) than what the pre-trained model was trained on. For all parts, my P@1 was 0, which I thought was strange.